

JBSHield: Defending Large Language Models from Jailbreak Attacks through Activated Concept Analysis and Manipulation

Shenyi Zhang and Yuchen Zhai, Key Laboratory of Aerospace Information Security and Trusted Computing, Ministry of Education, School of Cyber Science and Engineering, Wuhan University; Keyan Guo and Hongxin Hu, University at Buffalo; Shengnan Guo, Zheng Fang, and Lingchen Zhao, Key Laboratory of Aerospace Information Security and Trusted Computing, Ministry of Education, School of Cyber Science and Engineering, Wuhan University; Chao Shen, Xi'an Jiaotong University; Cong Wang, City University of Hong Kong; Qian Wang, Key Laboratory of Aerospace Information Security and Trusted Computing, Ministry of Education, School of Cyber Science and Engineering, Wuhan University

USENIX 2025

Gahyun Baek

JBShield


- Jailbreak attack: prompt designed to bypass the safety guardrails of an aligned LLM.
- Jailbreaks activate **compliance concepts**
not hide harmful intent, activate a specific jailbreak concept
- Toxic semantics remain recognized

JBShield

- Detection checks **dual activation**: toxic semantics & jailbreak concept
- Mitigation edits hidden states

Aligned LLMs Still Obey Jailbreaks


- Safety alignment refuses harmful prompts
 - Jailbreaks reframe unsafe requests
- => the same unsafe intent can lead to different model behavior depending on how the prompt is written.



Same unsafe request, different behavior

Prior Defenses Miss Causal Mechanisms

- Input filters catch surface cues
- Output filters stop interactions
- Prompt defenses add tokens
- Fine-tuning increases deployment cost
- Mechanisms remain underexplained



JBShield tries to solve this by analyzing the **model's internal representations** directly.

Research Questions

RQ1: Can aligned LLMs recognize the toxic semantics in jailbreak prompts ?


RQ2: How do jailbreaks change the outputs of LLMs from rejecting to complying ?



Defense mechanism of jailbreak

Linear Representations Enable Concept Analysis

- LRH maps concepts to subspaces
- Hidden states encode semantics
- Difference vectors isolate concepts
- Tokens interpret extracted directions



LRH: stating high-level concepts can be represented as directions or subspaces in hidden representations.

Toxic Concept = harmful/jailbreak prompt – benign prompt

Jailbreak Concept = jailbreak prompt – harmful prompt

JBShield Separates Two Activated Concepts

- **Toxic concept:** harmfulness
- **Jailbreak concept:** compliance
- Harmful prompts activate toxic only
- Jailbreaks activate both concepts
- Defense manipulates both directions



RQ1: LLMs Recognize Hidden Toxicity

LLMs can recognize toxic semantics in jailbreak prompts

- Harmful prompts yield warning tokens
- Jailbreaks yield similar toxicity
- Safety signal still exists

Harmful toxic tokens

caution · warning · disclaimer · ethical



Jailbreak toxic tokens


sorry · decode · translate · caution

Jailbreaks do not completely hide harmfulness from the model

RQ2: Jailbreaks Add Compliance Pressure

Jailbreak prompts add another concept

- Tokens include sure and yes
- Jailbreaks work by adding pressure toward cooperation, role-play, or decoding, which competes with the refusal behavior.



Concepts	Source Prompts	Associated Interpretable Tokens
Toxic Concepts	Harmful	caution, warning, disclaimer, ethical
	IJP	understood, received, Received, hell
	GCG	caution, warning, disclaimer , warn
	SAA	sure, Sure, sorry , assured
	AutoDAN	character, persona, caution, disclaimer
	PAIR	caution, warning, disclaimer, ethical
	DrAttack	caution, sorry, unfortunately, Sorry
	Puzzler	bekan, implement, pdata, erste
Jailbreak Concepts	Zulu	translate, sorry , transl, Translation
	Base64	decode, base, received, unfortunately
	IJP	understood , Hello, received , interpreted
	GCG	CHANT, Subject , plaat, bekan
	SAA	sure, Sure , mystery, CHANT
	AutoDAN	character, protagonist, persona , imagined
	PAIR	yes, sure, Sure, Subject
	DrAttack	sure, Sure , response, Response
	Puzzler	bekan, occas, CHANT, plaat
	Zulu	CHANT, translate , IMIT, translated
	Base64	decode , interpretation, received, reception

Mechanism Explains Refusal Failure

- Toxic signal triggers refusal
 - Jailbreak concept pushes compliance
 - Compliance can dominate toxicity
-
- **Defense should detect both toxic & jailbreak concepts**




JBShield Combines Detection and Mitigation

• JBShield-D

- detecting by checking whether both the toxic concept and the jailbreak concept are activated.



• JBShield-M

- mitigation by editing hidden states during generation.



JBShield-D Finds Critical Concept Layers

- Compute layerwise embedding gaps
- Select minimum cosine layer
- Separate toxic and jailbreak layers
- Calibrate anchor representations
- Store concept subspaces




(a) Jailbreak Detection: JBShield-D

*calibration: using a small set of benign, harmful, and jailbreak prompts to set the internal reference points for detection and mitigation


Detection Requires Both Concepts Activated

- Extract test toxic direction
- Compare with **toxic anchor**
- Extract jailbreak direction
- Compare with **jailbreak anchor**
- Flag when both exceed thresholds




$$R(x) = \begin{cases} \text{True}, & \text{if } s_t \geq T_t \text{ and } s_j \geq T_j, \\ \text{False}, & \text{else.} \end{cases} \quad (11)$$

Jailbreak = toxic && jailbreak




JBShield-M Edits Representations Directly

- Enhancing the toxic concept
 - Add toxic concept vector
- Weakening the jailbreak concept
 - Subtract jailbreak concept vector
- Scale from calibration projections
- Preserve normal functionality
- Generate safe content



(b) Jailbreak Mitigation: JBShield-M



Concept manipulation, not fixed refusal

Experiments Cover Diverse Attacks and Models

- 5 LLMs tested
- 9 jailbreak attacks evaluated
- 3 prompt datasets used
- 10 defenses compared
- Metrics: F1 (detection) and ASR (mitigation)

	IJP	GCG	SAA	AutoDAN	PAIR	DrA	Puzz	Zulu	B64
Mistral	Green	Blue	Green	Blue	Green	Blue	Green	Blue	Green
Vicuna-7B	Blue	Green	Blue	Green	Blue	Green	Blue	Green	Blue
Vicuna-13B	Green	Blue	Green	Blue	Green	Blue	Green	Blue	Green
Llama2	Blue	Green	Blue	Green	Blue	Green	Blue	Green	Blue
Llama3	Green	Blue	Green	Blue	Green	Blue	Green	Blue	Green

45 model-attack combinations

JBShield-D

- Average accuracy: 0.95
- Average F1: 0.94
- LlamaGuard F1: 0.75
- Works across five LLMs

Methods	Accuracy↑ / F1-Score↑								
	IIP	GCG	SAA	AutoDAN	PAIR	DrAttack	Puzzler	Zulu	Base64
Mistral-7B									
PAPI	0.04/0.08	0.05/0.09	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
PPL	0.01/0.03	0.33/0.48	0.00/0.00	0.00/0.00	0.01/0.01	0.00/0.00	0.00/0.00	0.95/0.95	0.00/0.00
LlamaG	0.68/0.81	0.78/0.87	0.83/0.90	0.77/0.87	0.74/0.85	0.84/0.91	0.77/0.87	0.50/0.67	0.58/0.73
Self-Ex	0.42/0.59	0.52/0.68	0.40/0.57	0.56/0.72	0.46/0.63	0.51/0.67	0.44/0.62	0.32/0.49	0.37/0.54
GradSafe	0.01/0.02	0.63/0.77	0.00/0.00	0.00/0.00	0.05/0.10	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
Ours	0.84/0.86	0.97/0.97	0.99/0.99	0.97/0.97	0.84/0.86	0.82/0.80	1.00/1.00	0.99/0.99	0.99/0.99
Vicuna-7B									
PAPI	0.04/0.08	0.14/0.25	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
PPL	0.01/0.03	0.47/0.62	0.00/0.00	0.01/0.02	0.00/0.00	0.00/0.00	0.00/0.00	0.95/0.95	0.00/0.00
LlamaG	0.65/0.79	0.75/0.86	0.85/0.91	0.72/0.83	0.75/0.85	0.84/0.91	0.75/0.86	0.49/0.65	0.55/0.71
Self-Ex	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.01/0.02	0.01/0.03
GradSafe	0.03/0.06	0.00/0.00	0.00/0.00	0.00/0.00	0.03/0.06	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
Ours	0.82/0.83	0.95/0.96	0.99/0.99	0.97/0.97	0.91/0.91	0.99/0.99	1.00/0.91	0.99/0.99	1.00/1.00
Vicuna-13B									
PAPI	0.04/0.08	0.02/0.04	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
PPL	0.01/0.03	0.79/0.86	0.00/0.00	0.01/0.02	0.01/0.02	0.00/0.00	0.00/0.00	0.95/0.95	0.00/0.00
LlamaG	0.64/0.77	0.76/0.86	0.84/0.91	0.75/0.76	0.76/0.86	0.85/0.92	0.75/0.85	0.48/0.64	0.54/0.70
Self-Ex	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
GradSafe	0.01/0.02	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
Ours	0.99/0.98	0.99/0.99	0.99/0.99	0.99/0.99	0.98/0.99	0.95/0.98	1.00/0.75	0.99/0.99	1.00/1.00
Llama2-7B									
PAPI	0.04/0.08	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
PPL	0.01/0.03	0.79/0.86	0.00/0.00	0.10/0.18	0.00/0.00	0.00/0.00	0.00/0.00	0.95/0.95	0.00/0.00
LlamaG	0.41/0.57	0.32/0.48	0.63/0.77	0.38/0.55	0.53/0.69	0.57/0.72	0.49/0.65	0.30/0.46	0.35/0.51
Self-Ex	0.31/0.33	0.28/0.32	0.36/0.39	0.27/0.31	0.27/0.30	0.32/0.35	0.24/0.27	0.30/0.33	0.29/0.32
GradSafe	0.39/0.56	0.97/0.98	0.00/0.00	0.96/0.98	0.62/0.77	0.00/0.00	0.18/0.31	0.00/0.00	0.00/0.00
Ours	0.84/0.86	0.82/0.86	0.93/0.94	0.98/0.98	0.87/0.88	0.99/0.99	0.81/0.85	0.91/0.91	0.92/0.93
Llama3-8B									
PAPI	0.04/0.08	0.02/0.04	0.00/0.00	0.02/0.04	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00	0.00/0.00
PPL	0.01/0.03	0.85/0.90	0.00/0.00	0.23/0.36	0.00/0.00	0.00/0.00	0.00/0.00	0.95/0.95	0.00/0.00
LlamaG	0.46/0.63	0.54/0.70	0.71/0.83	0.50/0.67	0.60/0.75	0.70/0.82	0.55/0.71	0.34/0.51	0.38/0.56
Self-Ex	0.15/0.26	0.12/0.21	0.19/0.31	0.11/0.19	0.16/0.26	0.16/0.27	0.18/0.30	0.12/0.21	0.14/0.24
GradSafe	0.41/0.58	0.21/0.35	0.00/0.00	0.97/0.98	0.37/0.54	0.00/0.00	0.92/0.96	0.00/0.00	0.00/0.00
Ours	0.91/0.92	0.98/0.99	1.00/1.00	0.97/0.97	0.77/0.86	0.97/0.96	0.99/0.99	0.99/0.99	0.97/0.97

0.95

Average accuracy

0.94

Average F1

JBShield-M

- No-defense ASR averages: 61%
- **JBShield** ASR: 2%

Models	Methods	Attack Success Rate↓									Average ASR↓
		IJP	GCG	SAA	AutoDAN	PAIR	DrAttack	Puzzler	Zulu	Base64	
Mistral-7B	No-def	0.56	0.92	0.98	1.00	0.82	0.74	1.00	0.48	0.40	0.77
	Self-Re	0.46	0.80	0.86	1.00	0.55	0.40	1.00	0.40	0.18	0.63
	PR	0.40	1.00	0.80	1.00	0.80	0.08	0.90	0.48	0.20	0.63
	ICD	0.52	0.45	0.58	1.00	0.70	0.68	1.00	0.06	0.08	0.56
	SD	0.52	0.70	0.96	0.98	0.78	0.86	1.00	0.32	0.40	0.72
	DRO	0.50	0.88	0.96	1.00	0.40	0.46	1.00	0.48	0.42	0.68
	Ours	0.24	0.36	0.12	0.00	0.08	0.04	0.00	0.02	0.00	0.10
Vicuna-7B	No-def	0.38	0.86	0.96	0.96	0.88	0.94	0.95	0.12	0.18	0.69
	Self-Re	0.34	1.00	0.88	1.00	0.70	0.62	0.95	0.18	0.00	0.63
	PR	0.22	1.00	0.82	1.00	0.75	0.34	0.80	0.40	0.22	0.62
	ICD	0.26	0.80	0.68	1.00	0.65	0.70	0.85	0.00	0.02	0.55
	SD	0.08	0.00	0.04	0.08	0.22	0.12	0.35	0.00	0.00	0.10
	DRO	0.36	1.00	0.64	1.00	0.60	0.52	0.95	0.54	0.06	0.63
	Ours	0.04	0.18	0.00	0.00	0.04	0.00	0.00	0.00	0.00	0.03
Vicuna-13B	No-def	0.36	0.78	0.92	1.00	0.68	0.98	0.95	0.0	0.10	0.64
	Self-Re	0.28	1.00	0.76	1.00	0.50	0.30	0.95	0.02	0.02	0.54
	PR	0.32	1.00	0.48	1.00	0.55	0.32	0.95	0.26	0.12	0.56
	ICD	0.28	0.75	0.52	1.00	0.70	0.78	0.45	0.00	0.02	0.50
	SD	0.04	0.02	0.02	0.02	0.08	0.00	0.00	0.00	0.00	0.02
	DRO	0.28	1.00	0.60	1.00	0.40	0.60	0.95	0.14	0.04	0.56
	Ours	0.00	0.00	0.00	0.00	0.00	0.02	0.00	0.00	0.00	0.00
Llama2-7B	No-def	0.26	0.50	0.60	0.60	0.30	0.32	0.95	0.14	0.30	0.44
	Self-Re	0.10	0.30	0.48	0.55	0.20	0.22	0.00	0.00	0.00	0.21
	PR	0.20	0.30	0.32	0.40	0.20	0.06	0.15	0.82	0.02	0.27
	ICD	0.02	0.25	0.36	0.70	0.05	0.12	0.00	0.00	0.00	0.17
	SD	0.32	0.00	0.00	0.00	0.24	0.10	0.40	0.00	0.42	0.16
	DRO	0.20	0.10	0.28	0.90	0.30	0.48	0.55	0.02	0.04	0.32
	Ours	0.02	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Llama3-8B	No-def	0.24	0.64	0.74	0.62	0.30	0.38	0.45	0.52	0.48	0.49
	Self-Re	0.02	0.15	0.44	0.30	0.05	0.36	0.00	0.02	0.00	0.15
	PR	0.26	0.10	0.14	0.10	0.20	0.04	0.05	0.46	0.06	0.16
	ICD	0.00	0.10	0.18	0.30	0.05	0.00	0.00	0.00	0.00	0.07
	SD	0.42	0.34	0.28	0.26	0.44	0.40	0.95	0.50	0.50	0.45
	DRO	0.24	0.20	0.42	0.50	0.10	0.12	0.00	0.60	0.14	0.26
	Ours	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.02	0.00	0.00

61%


Before defense

≈2%

After JBShield

Utility Degrades Less Than 2%

- MMLU performance largely preserved
 - Mitigation activates after detection
- > **Normal prompts mostly unaffected**



5-shot MMLU across five LLMs

Normal-input FPR averages 2%

My thoughts

Strength

- introducing “jailbreak concept”
- less affect to the utility

Weakness

- Requiring white-box access
- JBShield requires model-specific calibration
- they didn’t measure other attacks like multi-turn attacks

Thank you

Only 30 Samples Calibrate JBSHield

- N=30 balances performance
- Higher N may overfit
- Some attacks need N=10
- Calibration totals 90 prompts


* Interestingly, increasing the number of samples does not always improve performance.

Calibration Dataset Size <i>N</i>	Accuracy↑/F1-Score↑								
	IJP	GCG	SAA	AutoDAN	PAIR	DrAttack	Puzzler	Zulu	Base64
10	0.90/0.90	0.91/0.90	0.99/0.99	0.96/0.95	0.55/0.18	0.87/0.85	1.00/1.00	0.99/0.99	0.99/0.99
20	0.88/0.89	0.95/0.95	0.99/0.99	0.97/0.97	0.80/0.84	0.87/0.85	1.00/1.00	0.99/0.99	0.99/0.99
30	0.84/0.86	0.97/0.97	0.99/0.99	0.97/0.97	0.84/0.86	0.82/0.80	1.00/1.00	0.99/0.99	0.99/0.99
40	0.85/0.87	0.96/0.97	0.99/0.99	0.96/0.97	0.81/0.82	0.82/0.80	1.00/1.00	0.99/0.99	0.99/0.99
50	0.81/0.84	0.96/0.96	0.99/0.99	0.96/0.96	0.79/0.80	0.78/0.77	0.99/0.66	0.99/0.99	0.99/0.99

Average Mistral-7B F1 by N

Concept Extraction Converts Prompts Into Directions

- Form counterfactual prompt pairs
- Use last-token hidden states
- Compute representation differences
- Apply rank-one SVD
- Map direction to vocabulary



Conclusion

- Suggesting JBSHield with JBSHield-D and JBSHield-M.
- Jailbreak prompts work by activating a compliance-related concept while toxic semantics are still recognized.
- Based on this, it detects dual concept activation and mitigates attacks by editing hidden representations.

Jailbreak

- **Manually-designed**

- Human-crafted jailbreak prompts

Ex) IJP

- **Optimization-based**

- Use automated algorithms to exploit model gradients and generate malicious soft prompts

Ex) GCG, SAA

- **Template-based**

- Optimize adversarial templates that embed harmful requests

Ex) AutoDAN, PAIR

- **Linguistics-based**

- Hide malicious intent within seemingly benign inputs to bypass safety guardrails

Ex) DrAttack, Puzzler

- **Encoding-based**

- Transform or encode inputs to obscure harmful intent and evade LLM safeguards

Ex) Zulu, Base64